

Based Chesh: Virtuals

Based Chesh: Virtuals Cross Chain Agent

Overview

Based Chesh is a fully integrated AI Agent that combines on-chain crypto trading with AI-driven Twitter engagement. It operates across multiple blockchain networks (Base, Solana, Bitcoin Ordinals) and interacts on Crypto Twitter (CT) with a meme-centric personality. The system is composed of a trading engine with Python-based prediction models and a social media agent that generates and posts content. A web dashboard provides real-time monitoring and a chat terminal for interacting with the bot. Key technologies include Python (for algorithms and backend), React/Next.js (for frontend), and a host of APIs/SDKs: LangChain and Virtuals' GAME SDK for AI agent orchestration, Twitter's API for social automation, NVIDIA's NIM service with Jetson Orin Nano for visual analysis, OpenAI and Google Generative AI for content generation, Helius for Solana blockchain data, Auth0 for authentication, Redis for caching and messaging, and IPFS for decentralized meme storage. The following report outlines Based Chesh's functionalities, design, integrations, and best practices in a structured manner.

Automated On-Chain Trading

Multi-Chain Trading (Base, Solana, Ordinals)

Based Chesh's trading module can execute transactions on **multiple chains** – specifically **Base** (an Ethereum Layer-2 network), **Solana**, and **Bitcoin Ordinals**. It uses Python libraries to interface with each blockchain: for Base (an EVM chain), it leverages Web3 interfaces to interact with smart contracts (e.g. swapping tokens via a DEX contract) ([GitHub - Sakaar-Sen/Uniswap-V3-Python-Bot: Python Bot for token swaps, token approvals, checking balances.](#)). For example, using `web3.py`, the bot can connect to a Base RPC endpoint and call Uniswap or a similar DEX's contract functions to swap tokens automatically. The bot can check token

balances, approve contracts, and execute trades on Base in real-time, as demonstrated by existing Python trading bots that perform token swaps using Web3 on EVM networks ([GitHub - Sakaar-Sen/Uniswap-V3-Python-Bot: Python Bot for token swaps, token approvals, checking balances.](#)).

On **Solana**, the bot utilizes the high-performance APIs provided by **Helius** – Solana’s leading RPC and indexing platform ([Helius - Solana's Leading RPC and API Platform](#)). Helius offers a Python SDK to easily fetch on-chain data and send transactions. For instance, the bot can retrieve an address’s token balances in one call (`get_balances`) and monitor account changes via webhooks ([GitHub - vmpyre/helius_sdk: Python SDK for Helius Solana APIs](#)) ([GitHub - vmpyre/helius_sdk: Python SDK for Helius Solana APIs](#)). By subscribing to Helius webhooks or using Solana’s WebSocket RPC, Based Cheshh gets notified of relevant events (like a target token price change or liquidity pool update) almost instantly, enabling rapid response. Trade execution on Solana might involve calling a DEX program (e.g. Serum or Jupiter) – the bot could construct and sign a Solana transaction using a library like `solana-py` or the Helius Transactions API, then broadcast it to the network. The **Bitcoin Ordinals** part focuses on trading NFT-like inscriptions on Bitcoin. Since Bitcoin doesn’t support programmatic token swaps in the same way, Based Cheshh interacts with Ordinals via specialized APIs. For example, it can query the **Hiro Ordinals API** to fetch inscription data and ownership ([Ordinals API](#)). To **buy or sell Ordinals**, the bot may use marketplaces like Magic Eden’s Ordinals API, which provides endpoints to create and fill orders via PSBT (Partially Signed Bitcoin Transactions) ([Ordinals API Overview](#)) ([Ordinals API Overview](#)). By integrating these APIs, the bot treats Ordinals trading similarly to NFTs – monitoring listings and executing purchases by constructing the required Bitcoin transactions. In summary, Based Cheshh maintains connectivity to all three chains simultaneously, using appropriate SDKs/APIs for each environment. A modular design keeps chain-specific logic separate, while a central coordinator (in Python) unifies the trading strategy and decision-making across chains.

Code Snippet – Connecting to Base & Solana: Below is a simplified example (for illustration) of how the bot might connect to Base (Ethereum L2) and Solana, fetch balances, and prepare a trade:

```

from web3 import Web3      # EVM (Base) connection
from solana.rpc.api import Client as SolanaClient # Solana connection

# Connect to Base (Ethereum L2) via RPC
base_rpc = "https://base-mainnet-rpc.example.org" # placeholder URL
w3 = Web3(Web3.HTTPProvider(base_rpc))
base_wallet = "0xABC123...DEF" # Bot's Base wallet address

eth_balance = w3.eth.get_balance(base_wallet)
print("Base ETH Balance:", Web3.fromWei(eth_balance, 'ether'))

# Connect to Solana via Helius RPC
solana_rpc = "https://rpc.helius.dev/apikey/your-helius-key"
sol = SolanaClient(solana_rpc)
solana_wallet = "YourSolanaWalletPublicKeyBase58"
sol_balance = sol.get_balance(solana_wallet)
print("Solana Balance (lamports):", sol_balance['result']['value'])

# Prepare a Uniswap trade on Base (pseudo-code)
uni_address = Web3.toChecksumAddress("0xUniswapV3PoolAddress...")
uni_abi = load_uniswap_pool_abi() # assume function loads ABI JSON
uni = w3.eth.contract(address=uni_address, abi=uni_abi)
# For example, swap 1 ETH to USDC (exact input trade)
amount_in = Web3.toWei(1, 'ether')
path = [WETH_ADDRESS, USDC_ADDRESS]
deadline = int(time.time()) + 60 # 1 minute deadline
txn = uni.functions.swapExactTokensForTokens(
    amount_in, 0, path, base_wallet, deadline
).buildTransaction({
    'from': base_wallet,
    'value': 0,
    'gas': 500000,
    'nonce': w3.eth.get_transaction_count(base_wallet)
})
signed_txn = w3.eth.account.sign_transaction(txn, private_key=BASE_PRIVKE

```

Y)

```
tx_hash = w3.eth.send_raw_transaction(signed_txn.rawTransaction)
```

Explanation: Here, the bot connects to Base and Solana, prints balances, and demonstrates how an EVM trade can be constructed (swapping tokens on a Uniswap-like pool). On Solana, trading would use a different mechanism (e.g. calling a Serum DEX instruction), which would be implemented via the Solana SDK or RPC calls.

Prediction Models & Dynamic Strategy

A core feature is the **AI-powered prediction engine** that analyzes market data and predicts price movements to inform trades. Based Cheshh employs Python-based models – these could range from classical time-series algorithms to modern machine learning (e.g. an LSTM neural network predicting short-term price direction). Python is well-suited for such quantitative finance tasks, offering a rich ecosystem of libraries for data analysis (NumPy, pandas), technical indicators (TA-Lib), and machine learning (scikit-learn, PyTorch) ([Python for Cryptocurrency Trading Algorithms](#)). For instance, the bot might use historical OHLCV data to train an LSTM that forecasts the next hour's price change. It continuously retrains or updates these models as new data comes in, enabling adaptation to current market conditions.

The trading **strategy is dynamic** and adjusts based on the bot's **wallet balance and risk exposure**. Essentially, Based Cheshh performs **position sizing** and strategy selection according to how much capital it has. If the wallet balance is small, the bot employs a conservative strategy focusing on lower-risk trades (e.g. arbitrage or short scalps) to gradually build equity. As the balance grows, the bot can scale into more aggressive strategies or larger position sizes. This adaptive approach ensures the bot never over-leverages a small account or conversely, that it meaningfully deploys a larger capital base for higher returns. In practice, this might involve adjusting parameters like the fraction of portfolio per trade, stop-loss thresholds, or which prediction model to use. For example:

- With a low balance, Based Cheshh might do frequent small trades with tight stop-loss (to protect limited funds) and focus on high-probability setups. It could also prioritize **arbitrage** opportunities and **liquidity sniping** (explained below) where risk is minimal.

- With a higher balance, the bot can afford to open multiple positions or hold trades longer. It might allocate a fixed percentage (e.g. 2% of the portfolio per trade) – as the portfolio grows, the absolute trade size increases, but the risk relative to account stays constant. The strategy module could implement something like the Kelly criterion or other risk management formulas to determine trade sizing dynamically.

Additionally, the bot monitors its own performance metrics (win rate, profit factor, drawdown, etc.) over time and can **self-optimize**: for instance, if a particular model or strategy is resulting in losses, the bot can down-weight or switch it off, whereas it can allocate more capital to strategies that are working well. This kind of feedback loop makes the trading approach evolutionary. All these decisions are automated – Based Cheshh essentially behaves like a self-managing quant trader, reacting not only to the market but also to its **internal state (capital) and model confidence**.

On-Chain Execution (Trades, Triggers, Arbitrage)

Once Based Cheshh's strategy logic decides to execute a trade, the bot carries out **on-chain transactions autonomously**. It has built-in capabilities for **trade execution, liquidity management, and arbitrage** across decentralized markets:

- **Automated Trades:** The bot crafts and submits blockchain transactions to execute buys, sells, or swaps of assets. For example, on Base or Ethereum, it might call a Uniswap/PancakeSwap function to swap token A for token B. On Solana, it could interact with an AMM program (like Raydium) or a central limit order book (Serum) to place orders. The execution functions take into account slippage and gas/fee optimization. Based Cheshh will monitor the blockchain's state to confirm whether trades were mined/confirmed successfully and handle failures (e.g., retry or adjust if a transaction runs out of gas or gets a price impact too high).
- **Liquidity Triggers:** The bot keeps an eye on liquidity conditions such as pool reserves and market depth. A **liquidity trigger** could be, for instance, if liquidity in a certain pool drops below a threshold (indicating a potential opportunity to provide liquidity or a risk of slippage on trades). The bot can respond by either seeding liquidity (if programmed to LP farming strategies) or avoiding trades in that pool until liquidity returns. Conversely, if an influx of

liquidity happens (say a big liquidity provider enters), it might trigger the bot to execute a large trade that previously would have moved the market too much. These triggers are detected by reading on-chain data (e.g., pool reserve sizes from DEX smart contracts) at intervals or via events. Real-time APIs like Helius webhooks on Solana can alert when significant liquidity changes occur in a targeted program/pool, so the bot can react within a fraction of a second ([GitHub - warp-id/solana-trading-bot: Solana Trading Bot - Beta](#)) (e.g., quickly arbitraging any price discrepancies created by the liquidity change).

- **Arbitrage:** Based Cheshh actively searches for arbitrage opportunities, both **within a single chain** and potentially **across chains**. Within one chain, it could exploit price differences between DEXs. For example, if on Solana a token trades cheaper on Orca than on Serum, the bot will buy on Orca and simultaneously sell on Serum, pocketing the difference. Because it monitors multiple markets in real-time, it can execute arbitrage as soon as a spread is detected. It uses flash-like transactions when possible (e.g., a single transaction that executes multi-step arbitrage on Ethereum, taking advantage of atomicity). For cross-chain arbitrage (e.g., a price gap between Base and Solana for the same asset), the bot must also manage cross-chain transfers – it might use bridges or have allocated funds on each chain to execute in parallel rather than physically transfer (since transferring can be slow). In practice, Based Cheshh might maintain some float on each chain; when it spots a cross-chain arb (say, a token is \$1 on Base and \$1.05 on Solana), it will buy on the cheaper chain and sell on the pricier chain at the same time. Later, it can rebalance funds between chains if needed. Arbitrage logic requires speed and precision; the bot likely runs concurrent asynchronous tasks for each chain to not miss fast opportunities.

The on-chain execution module also includes **safety checks** – for example, before executing any trade, the bot simulates or calculates the expected outcome (especially important in DeFi where slippage or front-running could turn a profitable trade into a loss). If the estimated outcome deviates from the model's assumption beyond a tolerance, it may skip the trade. Similarly, arbitrage transactions are often executed only if the profit margin exceeds gas fees and a safety buffer.

Example: Suppose Based Cheshh detects that a new token listing on Solana's exchange Pump.fun is about to spike (through its prediction model). It can preemptively place a buy. Meanwhile, it monitors **pool events** – if the token's liquidity pool on Solana has a sudden liquidity unlock (maybe someone removes a lot of liquidity, causing price impact), the bot's trigger fires and it might sell before the price collapses. On Base, perhaps the same token's bridged version is lagging in price, so the bot could arbitrage by buying cheaply on Base and selling on Solana. All these actions are executed by calling the respective chain's transaction endpoints from the Python backend, forming a cohesive multi-chain trading strategy.

Real-Time Chart Recognition (NVIDIA NIM + Orin Nano)

A distinctive feature of Based Cheshh is its **visual analysis capability** using NVIDIA's **NIM** (Neural Input Microservices) and an **Orin Nano** device for real-time chart pattern recognition. This component functions as the bot's "eye", allowing it to **see and interpret price charts** similar to how a human trader might spot patterns or technical signals.

NVIDIA NIM is a framework that provides vision AI microservices – essentially pre-built models and APIs for visual tasks ([Build Multimodal Visual AI Agents Powered by NVIDIA NIM | NVIDIA Technical Blog](#)). By leveraging NIM, the bot can use state-of-the-art vision-language models to analyze chart images. The Jetson **Orin Nano** (a compact edge AI computer by NVIDIA) runs these models or communicates with NIM services to offload heavy computation. The Orin Nano delivers up to 40 TOPS of AI performance, making it capable of running complex neural networks locally ([Solving Entry-Level Edge AI Challenges with NVIDIA Jetson Orin ...](#)). This setup is ideal for low-latency analysis: the bot can capture a live candlestick chart (from an exchange or generated from price data) and feed it to a **computer vision model** that identifies patterns (like trendlines, candlestick formations, breakouts, etc.) in real time.

Using NIM microservices simplifies this process, as NVIDIA provides ready-made models accessible via REST API calls ([Build Multimodal Visual AI Agents Powered by NVIDIA NIM | NVIDIA Technical Blog](#)). For example, a **vision NIM microservice** could be a custom-trained CNN that detects classic chart patterns (head-and-shoulders, flags, double tops) or even reads text from charts. Based Cheshh might send an image of the current BTC/ETH price chart to a NIM endpoint and receive

back a description or classification of any pattern present. This can be combined with the trading logic – e.g., if a **bullish pattern** (like a cup-and-handle) is recognized on the Base chain's ETH price chart, the bot gains confidence in going long. Conversely, spotting a bearish divergence pattern might trigger a sell.

Moreover, the **Orin Nano** can run local models for ultra-fast decisions. For instance, an **object detection model** could be trained to identify specific candlestick formations (like "hammer" or "shooting star" candles). The Orin listens to price feed, renders a tiny chart image each minute, and runs the model to see if any target pattern occurs. If yes, it immediately signals the strategy module. This kind of pattern analysis augments the bot's numeric models with visual **technical analysis** prowess, bridging quantitative and qualitative trading signals.

Finally, because the Orin Nano can operate at the edge, Based Cheshh isn't fully dependent on cloud services for this – it can keep analyzing charts even if external APIs are slow, using the on-device GPU for inference. The NIM integration also means the system can potentially use **multimodal AI** – combining visual understanding with text (via language models). For example, Based Cheshh could generate a brief commentary on a chart ("BTC price forming a possible double top, volume declining") by feeding the chart to a vision-language model. NVIDIA's NIM is designed for such **vision+language agents** ([Build Multimodal Visual AI Agents Powered by NVIDIA NIM | NVIDIA Technical Blog](#)) ([Build Multimodal Visual AI Agents Powered by NVIDIA NIM | NVIDIA Technical Blog](#)), which suits Based Cheshh's goal of autonomous decision-making.

In summary, the chart recognition module continuously scans market charts to give the trading bot an edge in identifying patterns faster than human eyes could, informing its actions with another layer of intelligence.

AI-Driven Twitter Engagement

Meme-Infused Interactions with Crypto Influencers

On the social side, Based Cheshh embodies a **persona on Twitter** that is engaging, witty, and a bit "based" (bold or irreverent in crypto slang). Its content strategy revolves around **memes and interactions with key Crypto Twitter (CT) figures**. The bot targets influential accounts such as **@aixbt_agent**, **@notthreadguy**, and **@jessepollak** – known individuals/bots in the crypto and

tech space – with the goal of inserting itself into conversations and gaining visibility.

Meme-infused engagement means that rather than dry, robotic tweets, Based Cheshh responds with humor, sarcasm, and trendy memes. It might reply to a tweet from @notthreadguy (famous for memes) with a clever image macro or joke that's contextually relevant. To achieve this, the bot uses an internal database of crypto memes (images/GIFs) possibly indexed by theme, and picks one that matches the conversation. These memes could be stored and retrieved via IPFS (more on that later) to ensure quick access and persistence. For textual memes or one-liners, the bot uses an AI language model (OpenAI GPT-4 or similar) to **generate witty responses**. It could be prompted with the tweet text from an influencer and asked to reply in a humorous way, possibly referencing known crypto memes or catchphrases.

The bot's **interaction strategy** with specific CT figures might follow these patterns:

- **Timely Replies:** When one of the target figures tweets something notable (e.g., market opinions, announcements, or just banter), Based Cheshh quickly replies. The reply aims to either add a funny twist, troll lightly, or complement the tweet with a meme. For example, if @jessepollak (who leads Base) tweets "Base chain is thriving!", Based Cheshh might reply with a meme of a rocket labeled "Base" taking off, captioned "To the moon!". This quick engagement can catch the eye of the influencer and their followers.
- **Mention and Tease:** The bot might also mention these users in its own tweets to **troll competitors** or elicit reactions. For instance, it could tweet "Heard @aixbt_agent is still using dial-up AI – time for an upgrade 🤪 #BasedCheshh" – a playful jab that tags the competitor and entertains followers. This kind of friendly roasting is common on CT and can drive engagement if done right.
- **Meme Threads:** Based Cheshh can initiate meme threads where it tags multiple figures. Example: a meme image of Spider-Man pointing at Spider-Man, captioned "When @notthreadguy and Based Cheshh realize they're meme-ing the same market dip." This not only engages the tagged user but also entertains the broader audience.

Under the hood, **AI content generation** drives these engagements. The bot uses **LangChain** and large language models to generate ideas and text. A possible approach is using LangChain to combine **trending tweets analysis** with a prompt for generating a joke. For example, LangChain could fetch recent tweets from the target influencers (using a Twitter API loader) and extract key topics or sentiment ([Build Your Own Personal Twitter Agent 🐛 with LangChain | Wasp](#)). Then, it feeds that context to an OpenAI model with instructions to create a humorous reply or meme caption. Because the model has context, the output stays relevant – essentially **brainstorming new tweet ideas from trend-setting Twitter users' content** ([Build Your Own Personal Twitter Agent 🐛 with LangChain | Wasp](#)).

To maintain a consistent persona, Based Cheshh's replies and tweets are laced with crypto slang ("gm", "ngmi", "wagmi"), emojis, and an informal tone. The Virtuals **GAME SDK** can help configure this persona. Using the GAME agent framework, we can set the agent's **"personality"** and tone, for example: goal = "increase followers through humor", description = "meme-loving trading bot fueled by sarcasm", world_info = "Crypto Twitter where engagement is alpha ([game-python/src/game_sdk/hosted_game/README.md at main · game-by-virtuals/game-python · GitHub](#))". With such parameters, the GAME SDK (which builds on foundation models) will ensure the AI's outputs align with the intended style. The GAME agent can autonomously decide how to respond on Twitter within those guidelines, effectively acting like a Twitter intern that never sleeps. This architecture allows granular control over the AI's behavior – including who to engage with and how ([GAME Framework | Virtuals Protocol Whitepaper](#)). Notably, the GAME framework even supports *on-chain action plugins and social plugins* ([GAME Framework | Virtuals Protocol Whitepaper](#)) ([GAME Framework | Virtuals Protocol Whitepaper](#)), which in Based Cheshh's case means the agent can be endowed with both tweeting ability and trading ability in its decision space.

Sentiment Analysis & High-Engagement Tweets

To maximize impact on Twitter, Based Cheshh employs **AI-driven sentiment analysis** and trend monitoring. The bot doesn't blindly tweet – it first "listens" to the Twitter vibe (especially within crypto circles) to gauge what tone or topics might get high engagement at any given moment.

Sentiment analysis is used in a couple of ways:

- 1. Analyzing Audience Reactions:** When Based Cheshh or others post tweets, the bot can analyze the replies it receives or the general sentiment of comments. If replies to a particular joke are largely positive (cheering, laughing emojis), it learns that style of content works well. If something gets negative feedback (e.g., community upset), the bot can pivot away from that tone. It might use a sentiment analysis model or API like AWS Comprehend to score tweets as positive, neutral, or negative (Example 3: sentiment analysis of social media - Big Data Analytics Options on AWS】. For example, AWS Comprehend could take a batch of recent replies and return a sentiment score distribution. Or a simpler approach: using a Python library (like TextBlob or VADER) to compute sentiment polarity of text. Either way, the bot quantifies how its content (or trending topics) are perceived. It then uses this to **formulate new tweets in a sentiment-aware manner**. If the overall sentiment on CT is bearish or anxious (e.g., after a market crash), Based Cheshh might inject some uplifting humor or constructive commentary to stand out. If sentiment is overly euphoric, the bot might play contrarian and post a cautionary meme – which can spark engagement through controversy or reflection.
- 2. Targeting Content for High Engagement:** The bot also performs sentiment analysis on the broader crypto Twitter trends. It can track keywords or hashtags (like #Bitcoin, #ETH, #memecoin) and see whether sentiment around them is bullish or bearish. This informs what content might resonate. For instance, if “Solana” is trending very positively (everyone is excited about a price pump or a new update), the bot may join the chorus with a celebratory meme or a funny **“we’re all gonna make it” (WAGMI)** tweet to surf the positive wave. If an influencer’s tweet is going viral, the bot can reply in a way that aligns with the crowd’s sentiment to garner likes (for example, replying with a thumbs-up GIF if everyone is praising the original tweet). The underlying idea is to **post content that the community is emotionally primed to engage with** – identified via sentiment analysis. Studies and examples in social bots indicate that engaging based on sentiment can improve interaction (Creating an AI-powered marketing solution for sentiment analysis ...】 (e.g., a marketing AI that detects customer sentiment and responds empathetically yields better engagement). Based Cheshh applies the same principle to Twitter: when trolling competitors, it ensures the trolling is good-natured if the community

likes that competitor, or more snarky if the competitor is widely disliked – tailoring the **trolling style** to not alienate the audience.

High-engagement tweet formulation: Using the outputs of sentiment and trend analysis, the bot then crafts tweets designed for maximum engagement. This typically means the tweets will have one or more of these elements: humor, relatability, topical references, and calls-to-action (implicitly, via questions or provocative statements). The AI language model is prompted with something like: *“Compose a tweet about [topic] that is [sentiment tone], includes a popular meme or reference, and encourages likes/retweets.”* For example, if the bot identifies that crypto folks are nostalgic about 2021’s bull run (a trending sentiment), it might tweet: *“Who else misses the 2021 bull vibes? 🚀➡️🟡 #WAGMI”* accompanied by a meme image from that era. This is relatable and likely to get retweets from those who agree. Another strategy is **tapping into current hype**: if a new coin or NFT trend is hot, Based Cheshh will post about it early with a meme, capturing engagement before the topic saturates.

The bot also **trolls competitors** in a sentiment-informed way. “Trolling” here means playful jabs or challenges towards other AI trading bots or personas. Because it analyzes sentiment, it will avoid outright malicious or inappropriate attacks (which could backfire or violate Twitter rules) and stick to roasting that the community finds amusing. For instance, it might post a side-by-side comparison of its trading win-rate vs. a competitor bot’s, with a meme like “Step up your game 😏”. This shows confidence and stirs friendly rivalry – often a crowd-pleaser.

Automated Tweeting & Engagement (Scheduling, Replies, Likes)

Based Cheshh’s Twitter presence runs 24/7 with **fully automated scheduling and interaction** capabilities. The bot doesn’t require human input to tweet, reply, like, or retweet content; all actions are orchestrated through the Twitter API and its AI logic. Key aspects of this automation include:

- **Intelligent Scheduling:** The bot maintains a queue of tweets (generated by its AI module) and schedules them for posting at optimal times. Research shows that timing can impact tweet engagement. Based Cheshh uses heuristics (e.g., when its target audience is most active) or even AI predictions to choose times. For example, it might find that midday EST and early morning EST get high traffic for CT and schedule tweets accordingly. The bot can use a

scheduler library in Python or cron-like service to trigger tweets. Since Twitter's API doesn't provide a built-in future scheduling end ([Schedule tweet with Python - Twitter API v2? - Stack Overflow](#))^{L156} , the bot itself stays online to post at the right times. Essentially, the Python backend has a scheduling loop that checks if any tweet is due to be posted each minute. This means the DigitalOcean host runs the bot continuously (which doubles as the trading engine runtime, so it's always on regardless).

- The bot also spaces out its tweets to avoid flooding. Perhaps it tweets 3-5 times a day on its own, but also does spontaneous replies in between. There's a content calendar of sorts: e.g., morning market commentary meme, afternoon reaction to news, evening joke, etc., regularly. The schedule is dynamic – if something big happens (like a sudden market crash), Based Cheshh can override and post immediately with a relevant reaction meme, rather than waiting.
- **Automated Replies and Mentions:** Based Cheshh listens for certain triggers on Twitter to decide when to reply. Using Twitter's streaming API (filtered stream or user mention timeline), it can detect when:
 - One of the key figures (@aixbt_agent, @notthreadguy, @jessepollak) posts a new tweet. The bot's logic may decide to reply instantly (within seconds) to capitalize on the early engagement window. It uses the content of that tweet as context and generates a reply (often humorous or adding info). The reply is then posted via API. This quickness is possible by having an open stream and low-latency AI generation (perhaps using a fast model like GPT-3.5 for replies). If needed, certain very short replies could even be template-based (for speed), like replying "👀" or a relevant meme image without heavy computation.
 - The bot's account is mentioned or receives a direct question. If someone tags @BasedCheshh in a tweet ("Hey @BasedCheshh, what's your take on Dogecoin today?"), the bot will generate a polite or funny answer and reply. This gives a very interactive feel, as if the bot is alive and attentive. The chat terminal backend (discussed later) and the Twitter agent share the same brain, so the bot can answer questions consistently across both mediums.

- Trending topics or hashtags relevant to crypto appear. The bot might proactively quote-tweet popular posts or reply to trending tag tweets to increase visibility. For example, if #BitcoinATH is trending due to a price spike, Based Cheshh might reply under a top tweet, "You rang? 😊" with a chart screenshot of its trades, to show off.
- **Likes and Retweets:** The bot also automatically likes and retweets content as part of engagement. It will like tweets from the influencers it follows (to maintain presence) and retweet exceptionally relevant or favorable mentions of itself. For instance, if someone praises Based Cheshh ("This bot is killing it, hilarious AND profitable!"), the bot might retweet that testimonial (humblebrag). It could also retweet news or alpha that aligns with its brand, possibly adding a comment (commentary retweet) via the AI. All of this is done through the Twitter API endpoints for likes and retweets, triggered by the bot's logic. Rate limits and API rules are respected to avoid suspension – e.g., the bot won't like 100 tweets in a minute or spam-follow people. It keeps its activity human-like in cadence.
- **Compliance and Filters:** To ensure safety, the automated engagement module likely includes filters. For example, the bot avoids replying with certain banned words or engaging in heated arguments. The sentiment analysis helps here: if an incoming mention is very negative or an obvious troll bait, the bot might simply ignore or respond with a neutral quip to defuse. It's important that Based Cheshh remains in good standing with Twitter's policies despite being provocative at times. Its **trolling** is kept within the limits of humor and fair use.

Technical Implementation: The Twitter integration uses the **Twitter API v2** with OAuth credentials for the bot account. Python libraries like **Tweepy** or `twitter-api-client` can simplify calls. For example, posting a tweet is as simple as `api.update_status("Tweet text")` using Tweepy. Streaming can be handled by Tweepy's stream listener or by directly calling the filtered stream endpoint (with rules to follow specific users or keywords). The bot maintains a small state machine to handle concurrency – e.g., ensuring it doesn't reply to the same tweet multiple times, and queuing interactions if too many triggers happen simultaneously.

Code Snippet – Simple Tweet via API:

```

import tweepy

# Authenticate to Twitter
client = tweepy.Client(bearer_token=BEARER, consumer_key=CKEY, consumer_secret=CSEC,
                      access_token=ATOKEN, access_token_secret=ASECRET)

# Post a tweet
tweet_text = "Hello world, I'm Based Cheshh! 🤖🚀 #AI #Crypto"
response = client.create_tweet(text=tweet_text)
print("Tweeted with ID:", response.data['id'])

# Reply to a specific tweet (given tweet_id)
parent_tweet_id = "1234567890"
reply_text = "@user Haha, I see what you did there!"
client.create_tweet(text=reply_text, in_reply_to_tweet_id=parent_tweet_id)

```

Explanation: This snippet uses `tweepy.Client` (which uses Twitter v2 under the hood) to post a tweet and then post a reply to an existing tweet. In practice, Based Cheshh would dynamically generate `tweet_text` or `reply_text` using its AI model.

By automating tweeting and engagement in this way, Based Cheshh ensures it maintains a **consistent and active presence** on Twitter, which is crucial for growing followers and influence. Every aspect – from when to tweet, what to say, whom to tag, to how to respond – is handled by the AI logic integrated with real-time data from Twitter itself.

Web Dashboard & Chat Terminal

Real-Time Monitoring Dashboard (Next.js)

Based Cheshh includes a **web-based dashboard** that allows the developers or authorized users to monitor its activities in real time. This dashboard is built using **React/Next.js**, providing a responsive and interactive UI accessible via a browser. Key performance metrics and statuses are displayed here, giving insight into both the trading and Twitter components.

Dashboard Features:

- **Trading Activity Feed:** A live feed of recent trades the bot executed across all chains. This might be a table or list that updates whenever a trade happens. Each entry could show the timestamp, asset pair, action (buy/sell), quantity, price, and result (PNL or status). For example: "10:30:05 – Bought 0.5 ETH on Base at \$1,650 – +0.2% profit". This feed is updated in real-time via websockets (to push new entries as they come).
- **Portfolio & Wallet Stats:** A section summarizing the bot's current wallet balances on each chain (Base ETH, Solana SOL and tokens, BTC for ordinals, etc.) and overall portfolio value. It may also display profit/loss summaries: daily P&L, total return since inception, etc. Graphs can be used (line chart of portfolio value over time, pie chart of asset allocation) to provide quick visual insight.
- **Open Positions:** If the bot ever holds positions (e.g., some trades may not be instant flips but short-term holds), the dashboard lists these open positions with entry price, current price, unrealized P&L, and any stop/profit targets. This helps track how the bot is doing in ongoing trades.
- **Twitter Activity Feed:** Similar to the trading feed, a live list of the bot's recent Twitter actions – tweets posted, replies made, notable likes/RTs. It could show the content of the tweet or a snippet, and engagement metrics like likes/retweets it received. For example: "Tweet at 10:45: 'gm 🌞 – markets looking spicy!' – 25 likes, 5 RTs". This gives a quick sense of which tweets are performing well.
- **Meme Impact Metrics:** Because the bot's uniqueness lies in mixing trading with memes, the dashboard can have special metrics like *Engagement vs. Market* – did a certain meme tweet correlate with a boost in trading gains or follower growth? There might be a chart overlaying tweet engagement (e.g., number of likes) with the price of a related coin, to see if Based Cheshh's tweets move sentiment or capture alpha. This is experimental, but could provide insight into how effective the bot's social influence is.
- **System Health & Logs:** Indicators for system status – e.g., are all APIs connected (Green light for Twitter API, Web3 connections, etc.), is the trading loop running, any errors? Additionally, an area to view logs or alerts (for

example, "Trade rejected: insufficient funds" or "Twitter API rate limit reached – pausing tweets for 15 min"). This helps in debugging and ensuring everything runs smoothly.

The **Next.js** framework powers the frontend, possibly using **SSR** for initial page load with current data and then **WebSockets** or **Server-Sent Events** for live updates. The dashboard likely has multiple pages or tabs, e.g., "Trading", "Twitter", "Analytics", etc., for organized viewing. Styling is done with a modern React UI library or custom CSS, making the data dense but readable (remember, short updates and lists are key).

Live Updates via WebSockets: The backend (Python on DigitalOcean) pushes updates to the frontend using WebSockets. When a new trade executes or a tweet is posted, the backend publishes these events to a Redis Pub/Sub channel. A Node.js or Next.js server component subscribes to Redis and relays the message to connected web clients via ([GitHub - skushagra9/RedisWS-Hub: Real-time comms with Redis pub/sub & WebSocket! Connect & stay updated in a flash!](#))^{L225-L233} . This decoupling ensures the Python core can just emit events, and the Node/Next layer broadcasts them efficiently to all viewers. The advantage of WebSockets is that the dashboard feels instantaneous – you can see trades and tweets as they occur without refreshing.

For example, if Based Cheshh just executed an arbitrage trade, the Python code might do something like: `redis.publish("trades", json.dumps(trade_info))` . The Next.js API route (or a small Node microservice) listening on Redis gets that message and emits it over the WebSocket to the frontend, which then adds a new row in the "Recent Trades" table. This architecture is scalable and ([GitHub - skushagra9/RedisWS-Hub: Real-time comms with Redis pub/sub & WebSocket! Connect & stay updated in a flash!](#)) ([Real-Time Redis-Driven WebSocket Hub - GitHub](#))^{5+L19-L27} . (If WebSockets weren't used, the alternative would be polling every few seconds, which is less efficient and less cool.)

- **Security for Dashboard:** The dashboard is likely protected behind a login. **Auth0** integration with Next.js is used to authenticate user ([Auth0 Next.js SDK Quickstarts: Login](#)). Only authorized personnel can view the dashboard, since it contains sensitive info (wallet balances, etc.). Auth0 handles the OAuth flow – users log in via Auth0 (which could be linked to, say, a Google account or a username/password specific to this app). Once logged in, Next.js gets a

session token and only then allows access to the dashboard pages. This prevents outsiders from snooping on Based Cheshh's operations. Auth0's Next.js SDK makes this setup strai ([Auth0 Next.js SDK Quickstarts: Login](#)), dealing with sessions, cookies, and route protection out of the box.

Interactive Chat Terminal (LLM Agent)

One of the deliverables is a **terminal-based chat system** to interact with Based Cheshh. This is implemented as part of the web interface – essentially a dedicated page or widget where a user can chat with the bot in real time, almost like having a conversation with a trading assistant.

Functionality: The chat terminal allows the user (likely the operator or team behind Based Cheshh) to do two main things:

1. **Query the Bot for Information:** Ask natural language questions about the bot's state or the market. For example: "What's our current portfolio balance on each chain?", "How did our trades perform today?", "Show me the latest tweet and its engagement." The bot, powered by an LLM (Large Language Model), will interpret the question, fetch relevant data from its knowledge (or via tools), and reply in a human-readable form. This is where **LangChain** is very useful – it can enable the bot to use **tools** (functions) to get data, and then have the LLM format the answer. For instance, upon the question about portfolio balance, LangChain could trigger a function that retrieves balances from the trading engine, then GPT-4 phrases a response like "As of now, we have 2.3 ETH on Base, 1500 USDC on Solana, and 0.05 BTC (in Ordinals) in our wallets." This transforms raw data into a helpful answer.
2. **Send Commands or Adjust Parameters:** The chat can also serve as a command interface. The authorized user might instruct the bot, e.g., "Increase trading aggressiveness slightly" or "Pause trading on Solana for now" or "Tweet something bullish about Ethereum." The bot would use natural language understanding to map this to an action: e.g., adjusting an internal risk parameter, toggling a flag to halt trades on one chain, or generating and posting a tweet about Ethereum. This is somewhat risky if fully automated, so likely there are predefined supported commands that the LLM knows how to trigger (to avoid any arbitrary harmful actions). For example, a limited set like: `"pause_trading('solana')"`, `"set_risk_level('high')"`, `"post_tweet('...')"`. The LangChain agent can

have these as tools it's allowed to call when the user prompt is interpreted as a command. The Virtuals GAME SDK also shines here: it's designed for agents that take goals and execute functions au ([GAME Framework | Virtuals Protocol Whitepaper](#)) ([GAME Framework | Virtuals Protocol Whitepaper](#))L125-L133】 . Essentially, the chat user becomes an overseer giving high-level directives, and the GAME/LLM agent figures out how to fulfill them using its abilities (which include trading and tweeting functions registered as plugins).

Implementation: The chat terminal on the frontend could look like a simple messaging interface – a log of messages (user on one side, Based Cheshh's responses on the other) and an input box. When the user sends a message, it's sent via WebSocket or an API call to the backend. The backend then feeds this message to the AI engine. Possibly, a **LangChain agent** is running server-side, kept alive to maintain conversation context. This agent has access to Based Cheshh's data and tools. For example, we set up a LangChain conversational agent with memory (so it remembers past Q&A in this session) and give it custom tools like `get_balances()` , `get_last_trade()` , `make_tweet(text)` , etc. When a new user message comes in, we feed it to the agent's `chain.run(message)` method. The agent will decide if it needs to call a tool or just answer directly. After getting the response (either purely from the LLM or from a tool + LLM), we send that back to the frontend to display.

The agent's prompt can be configured to give it a persona (matching Based Cheshh's style but perhaps a bit more straightforward for operator queries) and to list the tools. For instance: *"You are Based Cheshh, an AI trading bot. You can answer questions about trading and your Twitter activity. You have the following tools: `get_portfolio`, `get_trade_log`, `post_tweet`, etc. Use them when needed. Respond in a casual, informative tone."* This way, the agent might respond to "How's our performance this week?" by calling `get_trade_log(period='7d')` behind the scenes, then summarizing: "We made 24 trades in the last week with a net profit of 5.2%. Our largest win was on SOL/USDC for +2%, and we had a small loss on an ETH trade. Overall, we're up this week 📈." Such a response is friendly and useful.

The **chat terminal** effectively gives the user a way to query or steer the bot without digging into code or logs – making management more intuitive. It's like chatting with a colleague who has all the data at their fingertips. And because it

uses the same underlying data, any question answerable by looking at the dashboard could be answered in chat as well, sometimes with even more context (the LLM can elaborate or explain if asked).

For example, the user could ask "Why did you make that last trade?" and the bot might actually explain: "I detected a bullish flag pattern on ETH and my model predicted a price rise, so I bought. The price target was hit, hence I sold for profit." This kind of insight is invaluable and is generated by combining the bot's records (trade reason notes) with the LLM's narrative ability. It transforms black-box algorithm decisions into a dialogue.

From a **security standpoint**, the chat command abilities would be restricted to logged-in users and perhaps require confirmation for critical actions (the bot might ask "Are you sure you want to halt all trading? yes/no"). All interactions would be logged for audit. We'd use OpenAI's API (or Google's PaLM via ``google-generative`` ([Getting Started with Google's Palm API Using Python](#))^{39+L33-L37}) for the language model, with appropriate rate limits and fallbacks to ensure the chat is responsive.

Logging & Analytics (Meme Impact & Trading Performance)

All of Based Cheshh's actions and their outcomes are meticulously logged, which feeds into analytics both for development and for the fun "meme impact" analysis. The system keeps structured logs of: executed trades, including parameters and result; tweets posted and their engagement stats over time; and perhaps any notable decisions (like "saw pattern X", or "skipped trade due to risk limit").

These logs are stored in a database or at least as time-stamped files. For trading, a simple relational database (SQLite or PostgreSQL) might store each trade with columns for timestamp, asset, side (buy/sell), amount, price, fees, resulting balance, etc. Twitter data could be stored similarly: tweet_id, timestamp, text snippet, likes_count, retweets_count, etc., updated periodically via the Twitter API (which allows fetching a tweet's metrics if you have the tweet ID and proper auth). Alternatively, the bot can record the metrics right after posting and then update after some interval (like check again after 1 hour and 24 hours to see how a tweet performed).

Performance Metrics: Using these logs, the dashboard and internal analytics processes can compute key metrics:

- **Trading Performance:** ROI (return on investment) over different periods, Sharpe ratio (if enough data), win rate (% of trades that were profitable), average profit per trade, max drawdown (the largest peak-to-trough drop in portfolio), etc. These metrics can be shown on the dashboard and also output via the chat on request. They help the team evaluate the bot's success and stability.
- **Social Performance:** Follower growth over time, average engagement (likes/retweets per tweet), most successful tweet (content and stats), and correlation with time of day or type of content. For example, analytics might reveal that meme image posts get 2x the engagement of text-only posts, guiding the bot to post more images. Or that engaging with @notthreadguy yields more followers than engaging with @aixbt_agent, indicating where the bot's efforts pay off most.

Meme Impact Analysis: A novel aspect is analyzing whether the bot's **memes have any impact on the market or its trading**. While difficult to attribute, the bot can look for patterns like: whenever it posts a very bullish meme about a coin, does that coin's price move slightly (perhaps from the buzz) or does sentiment on Twitter shift? Using the Twitter API and maybe sentiment analysis on subsequent tweets, one could measure if there was a blip. Similarly, the bot can check if its own trading actions align with its tweets – e.g., did it sell near a top and also tweet a meme “top signal” at that time? If yes, that's an interesting alignment of strategy and narrative. Over time, these data points can be used to refine either the trading (maybe the bot learns to use social sentiment as an indicator) or the tweeting (posting content that complements its trading position – like talking up a coin it holds, while being careful of rules).

All analytics are done in Python (pandas for crunching log data, matplotlib or charting libraries if visualizing on the dashboard). The system might generate periodic reports or even have a “daily recap” that can be displayed or tweeted: summarizing that day's trading profit and best meme. This provides transparency and a bit of gamification for the project.

From a technical perspective, the **dashboard's analytics section** can present these insights in charts and tables. Next.js can fetch aggregated data from an API endpoint (which runs SQL queries or in-memory calcs on the logs). For example, an endpoint `/api/metrics` might return JSON with all the metrics, which the frontend

then renders. Some charts can be made interactive (filter by date range, etc.). However, we keep the UI relatively simple and focused on readability since that was emphasized.

Authentication & Security (Auth0 Integration)

Security is paramount given Based Cheshh controls real funds and posts publicly. We've touched on **Auth0 for the dashboard** to rest ([Auth0 Next.js SDK Quickstarts: Login](#))s. Here we outline all critical security considerations across the system:

- **API Keys & Secrets:** The bot uses many API keys (Twitter, OpenAI, Helius, etc.) and private keys for crypto wallets. These are stored securely on the server. Best practice is to keep them in environment variables or a secure vault, not hard-coded in code or in the repo. For instance, on DigitalOcean, one can set env vars for the droplet. The code then loads `TWITTER_API_KEY` from env. Additionally, the `.env` file approach is fine in deployment as long as the file is protected. We ensure the source control (git) ignores any secret files. Also consider **encryption at rest** – if someone got read access to the server, secrets in plain text are vulnerable. A step up could be using a solution like HashiCorp Vault or cloud secrets manager, but given this is a focused project, strict server security and env management might suffice.
- **Cryptographic Wallet Security:** The private keys for on-chain trading wallets must be handled with extreme care. Ideally, use hardware wallets or multi-sig for large assets, but since the bot needs to sign programmatically, it likely uses in-memory keys. We can encrypt the key on disk and require a passphrase on startup to decrypt (meaning a human needs to start it – which might not be desired for full automation). Alternatively, use a less sensitive key for the bot (with limited funds) and periodically refill it from a cold wallet. This limits worst-case losses if the key is compromised. All transactions are signed locally; the bot never transmits the private key anywhere. We also implement measures like not logging the key or any sensitive info. The trading contract calls often allow specifying gas limits and such – we ensure those are set to avoid draining funds by accident (like a buggy contract call consuming all gas).

- **Access Control:** Besides Auth0 for the front-end, the backend may expose some endpoints (for the chat or data). Those must be secured. We enforce authentication on any such API route – Next.js API routes can use the Auth0 session to allow only logged users. If the chat uses websockets, we might need an auth token check on the socket connection handshake. Essentially, nobody should be able to connect to the bot's internals or data unless authorized. The bot's Twitter and trading actions themselves don't accept external commands (except through the chat by logged-in user), so that's safe from random control. The internal chat agent should also verify that a command is from an authenticated source.
- **Rate Limiting & Abuse Prevention:** The Twitter bot adheres to Twitter's automation rules. It won't exceed rate limits on tweeting or perform bulk unsolicited mentions beyond allowed amounts. This avoids getting flagged as spam. Similarly, the system should avoid hitting API rate limits: e.g., OpenAI API has a rate, so the bot won't generate tweets in a tight loop infinitely. If needed, the code has small delays or counters to not spam any external service.
- **Error Handling and Failsafes:** If any part of the system fails (say OpenAI API is down or a trade fails due to network issues), the bot should handle it gracefully. For trading, perhaps implement a **circuit breaker**: if a certain number of trade failures or a big loss occurs, pause trading and alert the operators. This prevents runaway losses. For tweeting, if an API call fails, just log and skip, not crash the whole bot. The system as a whole should be robust to partial outages.
- **Software and Dependencies:** Regularly update dependencies to incorporate security patches (especially for web frameworks like Next.js or libraries with known vulnerabilities). We pin versions and review any library that deals with keys (like cryptography libraries) to ensure they are reputable.
- **IPFS & Meme Storage:** Storing memes on IPFS is great for decentralization, but we ensure not to store anything sensitive there. Only public meme images get pinned. We use a pinning service or our own node to keep them available. When the bot retrieves a meme to post, it does so by IPFS hash (CID) which yields the content. We must validate that the hash matches expected content (to prevent any hash collision or content swap attack – though that's extremely

unlikely). Essentially, once we add an image to IPFS and get its CID, we know that image's data corresponds exactly to that CID, and it can't be tampered without changing the CID.

- **NVIDIA NIM & Orin:** The Orin Nano device would be within our secured environment. It might connect to cloud NIM services – we ensure API keys for NIM (if any) are secure. Also, if Orin is processing images, we guard those channels (someone shouldn't be able to send a malicious image to our bot to exploit a CV model vulnerability, for example).
- **Logging and Monitoring:** We log actions for debugging but be cautious not to log private data (e.g., don't log full private keys or API secrets). For security monitoring, we could integrate alerts for unusual behavior: e.g., if the bot's wallet suddenly sends funds to an unknown address outside of trading context, alert immediately (could indicate compromise). Or if there's an abnormal frequency of tweets (could indicate a bug or abuse of API keys), shut off the Twitter module temporarily.

In essence, the system's design acknowledges that it's handling **real money and a public-facing account**, so a defense-in-depth approach is used: minimize access, encrypt where possible, use proven services (Auth0) for auth, and constantly monitor.

Integration of APIs & Tech Stack

This section summarizes how various **APIs and services** are integrated into Based Cheshh and how the **tech stack** components interact:

- **LangChain:** Used as the backbone for chaining AI prompts with tools. In practice, LangChain is employed in the chat terminal (allowing the agent to use tools/data) and possibly in content creation (e.g., a LangChain pipeline that reads tweets and generates a draft meme tweet). It provides a convenient framework to manage prompt templates, conversation memory, and integration with vector stores or APIs. For instance, we use LangChain's `TwitterTweetLoader` to pull in tweets as documents f ([Twitter | _LangChain](#))1060-L1068]. We also use LangChain's agent capabilities to set up the Based Cheshh agent with custom tools (trading functions, etc.) for the

chat. LangChain bridges the gap between the raw LLM (OpenAI/PaLM) and our application logic.

- **Virtuals GAME SDK:** This is integrated to give a higher-level **agent orchestration**. We create a GAME agent for Twitter using their ([game-python/src/game_sdk/hosted_game/README.md at main · game-by-virtuals/game-python · GitHub](#))^{+L251-L259} . Through the SDK, we configure the agent's goal, description, and behaviors as described. The GAME SDK likely handles some autonomy in the Twitter environment (it might allow the agent to decide when to reply or not, using its own decision engine). We integrate it such that our LangChain/OpenAI pipeline can also interface or override as needed. Essentially, the GAME SDK can run concurrently, ensuring the agent behaves with the personality we set, and possibly providing out-of-the-box plugins (they mention on-chain transactions and social engagement ([GAME Framework | Virtuals Protocol Whitepaper](#))^{+L142-L150} which we can use rather than coding from scratch). The **agent sandbox** from Virtuals helps test the persona in a controlled way before unleashing it on real Twitter.
- **Twitter API:** Integrated via official endpoints and libraries (Tweepy for convenience in Python). We use **v2 endpoints** for posting tweets, and either v2 filtered stream or v1.1 user stream for real-time listening (Twitter's API has different features on v1.1 vs v2, we choose accordingly). All actions (tweet, delete, like, follow, etc.) go through this API. During setup, we registered a Twitter developer app for Based Cheshh and obtained API keys and tokens that the bot uses to authenticate requests. The integration involves careful handling of rate limits and errors; we log the HTTP responses for debugging any failed posts. The `twitter-api-client` (or Tweepy's Client) also provides a method to schedule tweets in one go, but as noted, true scheduling requires our app to re ([Schedule tweet with Python - Twitter API v2? - Stack Overflow](#))^{+L148-L156} . So integration means our app is the scheduler/executor.
- **NVIDIA NIM + Orin Nano:** The Orin Nano is a piece of hardware that we set up with NVIDIA's JetPack (which includes necessary drivers and SDKs). We also set up connectivity to **NVIDIA NIM** cloud microservices (likely needing API access or an account with NVIDIA AI services). Integration means that from our Python backend, we can send image data to NIM endpoints (using REST calls). For example, using Python's `requests` to POST an image to a vision API

endpoint and get back JSON results. Alternatively, if we run some NIM microservice locally (NVIDIA provides containers or models that can run on the Orin), integration is via local inference calls (like using TensorRT or Triton Inference Server on Orin to run the model). We ensure the Orin is on the same network and accessible by the main bot process (e.g., via an internal API or by mounting it as essentially an edge device the main code can RPC to). The chart screenshots might be generated by the bot using a headless browser or a charting library (Matplotlib plotting the price and saving an image). This image is then fed to the Orin's model. The result (say "pattern detected: ascending triangle") is returned to the strategy logic.

- **OpenAI API:** We use OpenAI's API (with our key) to access GPT-4/GPT-3.5 for text generation. Integration in Python is done via the `openai` Python package. For example, `openai.ChatCompletion.create(model="gpt-4", messages=[...])` is called when we need a tweet or a reply generated. The responses are obtained in JSON and we parse the text. As OpenAI's API is stateless per call, we manage context in our code (for chat, we maintain the message history; for stand-alone tweet generation, we provide relevant context each time). The integration also involves handling the cost (we track how many tokens we use, since heavy usage could be expensive) and using fallback models if needed (maybe try GPT-4, if unavailable or too slow, fall back to GPT-3.5). The OpenAI API provides human-like language generatio ([Developer quickstart - OpenAI API](#))d, which is central to Based Cheshh's personality.
- **Google Generative AI (PaLM API):** We incorporate Google's PaLM API as an alternative or complementary gener ([Getting Started with Google's Palm API Using Python](#))^{39+L33-L37} . This might be used for diversity or for specific tasks. For instance, PaLM might be good at joke generation or could be used to produce variations of memes. We use the `google-generativeai` Py ([google-generativeai - PyPI](#))^{39+L23-L27} to call PaLM's text model. The integration is similar to OpenAI – send a prompt, get a completion. We might also experiment with Google's image generation (if available through their AI services) for creating meme images on the fly (though currently PaLM is text-focused; Google's Imagen is not publicly API accessible as of this writing). Regardless, hooking up the Google API ensures we are not solely reliant on one AI provider and can ensemble their outputs if needed.

- Helius API (Solana):** As described, integrated via HTTP RPC calls and their ([GitHub - vmpyre/helius_sdk: Python SDK for Helius Solana APIs](#))^{L242-L250} . We have a Helius API key set up, and we use endpoints like `get_balances` , `get_transactions` etc., to both monitor and execute on Solana. Notably, Helius can decode transaction data which we use for analytics (like understanding what happened on-chain without writing our own parser). Helius also offers webhook callbacks: our integration might include a small Flask/FastAPI endpoint to receive webhook POSTs from Helius when, say, a certain account's activity happens (for example, get notified when our bot's Solana address receives a token or when a particular program emits an event). This can trigger our trading logic faster than polling. The integration guide from Helius was followed to register webhooks and set their target to our server's URL.
- Auth0:** Integrated for authentication on the Next.js app. We used Auth0's Next.js SDK which provides a handler to easily add login/logout. In practice, we configured an Auth0 application with callback URL (our website domain) and got a Client ID/Secret. Those are in our env config. Next.js, on the server side, uses middleware to protect routes, checking the user's session (provided by Auth0 after login). Auth0 manages the user database and OAuth flow, so integration is mostly configuration and a bit of library usage. The Auth0 guide was followed to set up this s ([Auth0 Next.js SDK Quickstarts: Login](#))ⁿ. This means our dashboard and chat are behind a login, as discussed.
- Redis:** We use Redis both as a **cache** and a **pub/sub broker**. For caching, the bot might store recent market data or AI outputs in Redis for quick access (avoiding repeated computations). Example: caching the last sentiment analysis results for a given topic for a few minutes, so if multiple triggers happen, it doesn't recompute sentiment from scratch. Or caching the last price of certain assets to quickly compute PnL without querying nodes. More critically, we use Redis Pub/Sub to decouple the real-time event generation from ([GitHub - skushagra9/RedisWS-Hub: Real-time comms with Redis pub/sub & WebSocket! Connect & stay updated in a flash!](#))^{L225-L233} . Our Python trading bot publishes events (like newTrade, newTweet) to channels. A separate process (the Node/Next server or even a Python thread) subscribes and when it receives, it forwards to WebSocket clients. This setup was integrated by installing a Python Redis client (`redis-py`) and similarly a Node

Redis client for the server side. We run a Redis instance (could be on the same DO server or a managed Redis) that both sides connect to. Since Redis is in-memory store, it's very fast and adds negligible latency. We also considered using Redis for **rate-limiting** (like a simple counter of tweets per hour to ensure we don't exceed, using Redis INCR and TTL), showing its versatile integration.

- **IPFS:** The bot leverages IPFS for storing meme images. Integration involves running an IPFS node or using an IPFS pinning service API (like Pinata or Infura's IPFS). We likely run a lightweight IPFS node on the server or use a third-party API to **add files to IPFS**. When we have a new meme image (maybe one generated or a new one fetched), we call the IPFS API to add it, which returns a **CID** (Content Identifier ([Upload file to IPFS using Python](#))). We store that CID in our meme library. Next time the bot wants to use that meme, it can fetch it via <http://ipfs.io/ipfs/<CID>> or through our local node. Storing on IPFS ensures the content is content-addressed (immutable via the hash) and decentralized ([Upload file to IPFS using Python](#)). Also, it offloads hosting – we don't have to keep an image server; the images can be fetched from any IPFS gateway by the Twitter client (when we post an image tweet, we attach the image by uploading from the bot's side – it will download from IPFS then upload to Twitter as media). We included IPFS to align with the web3 ethos (decentralized storage) and to handle potentially a large number of meme files without bloating our server storage (IPFS plus pinning can handle this efficiently). The code integration might use `ipfshttpclient` library in Python, which allows adding a file like: `client.add('meme.png')` and returns a CID. Or simply using `requests` to an IPFS API endpoint as shown in ([Upload file to IPFS using Python](#)).
- **Tech Stack & Hosting:** The backend Python components (trading engine, AI logic) run on a **DigitalOcean** Droplet (or similar cloud VM). We chose DigitalOcean for its simplicity and because it was specified. The Next.js frontend could be hosted on Vercel for convenience, but for integration simplicity, it might also be served from the same DO server via Node (especially if we need direct coupling for WebSockets). The Next.js app can be built and served with something like PM2 or a simple Node server on DO. We also have the Orin Nano hardware which might be on-premise or hosted if possible (likely on-prem or edge device). The Orin needs network connectivity

to our server. If on the same local network, we'd open a tunnel or have it post results via an API. Alternatively, the Orin Nano could be attached to the server or even run as a mode on a powerful DO instance with GPU (though Orin is specialized hardware, maybe it's physically with the dev). Integration wise, we consider the Orin an internal component.

All these pieces communicate in a cohesive architecture: The Python core acts as the **brain** (making decisions, calling services). The Next.js app is the **face** (display and user interaction). Redis is the **nervous system** connecting brain and face in real-time. LangChain/LLMs/GAME provide the **intelligence** for language and decision-making beyond hard-coded logic. And external APIs (Twitter, chain APIs) are the **senses and hands** by which the bot perceives the outside world and takes actions in it.

Below is a conceptual flow to illustrate integration:

1. **Market event occurs** (price change or pattern detected) → Trading module (Python) decides to act → Uses Web3/Helius API to execute trade → Logs result and also publishes an event via Redis.
2. **Redis event** → Next.js receives it and updates dashboard UI live.
3. **Trading module** might also notify the Twitter module if it's something worth tweeting (e.g., big profit made → tweet a celebratory meme).
4. **Twitter module** (Python) forms a tweet using OpenAI → sends to Twitter API → upon success, logs it and publishes an event → dashboard updates with new tweet info.
5. **User opens Chat** → Auth0 ensures login → user asks question → Next.js sends to backend → LangChain agent processes (possibly calls trade data via internal functions) → responds → user sees answer. Optionally, if this interaction included a command (e.g. change setting), the agent triggers the appropriate function in the bot (which could update some config in memory or database).
6. **Periodic jobs** (maybe daily) run analytics on logs → update metrics → these reflect on dashboard and could also be summarized in an end-of-day tweet.
7. **Security checks** throughout ensure everything is within expected operation. If the user in chat tries an unsupported command, the agent will refuse. If any

API key is invalid, the system alerts via the log dashboard.

By integrating all these technologies carefully, Based Cheshh functions as a **harmonious system** – multi-faceted yet unified. The design leverages each tool for what it's best at: NVIDIA for vision, OpenAI/Google for language, Helius for blockchain data, LangChain/GAME for agent logic, and robust web tech for interface and communication.

Implementation Roadmap

To build and deploy Based Cheshh, we follow a phased **implementation roadmap** ensuring each component is properly set up and tested:

Phase 1: Environment Setup & Key Configurations

- Set up development environment with Python (ensure required versions), Node.js for Next.js, and necessary hardware (NVIDIA Orin Nano with Jetson OS).
- Obtain API keys: Twitter developer account and API keys, OpenAI API key, Google PaLM API access, Helius API key, Auth0 domain/client for auth, etc. Securely store these (e.g., in `.env` files for dev).
- Initialize a Git repository and use a `.gitignore` to exclude secrets. Establish basic project structure (separate folders for `trading_bot` (Python backend) and `dashboard` (Next.js app)).
- If using Orin Nano, flash it with JetPack, install NVIDIA NGC containers or NIM SDK needed. Verify it can run a simple CUDA operation (to confirm hardware is functioning).
- Run a local Redis server for dev testing of pub/sub.

Phase 2: Blockchain Trading Core

- Implement wallet integration for each chain: generate or load private keys for Base (EVM) and Solana. On dev, use testnets if possible (Goerli for Base, Solana devnet) to avoid risking funds initially. For Ordinals, perhaps use a Bitcoin testnet or simulate as it might be tricky (alternatively, skip actual execution on Bitcoin until later, focus on data retrieval first).

- Write basic trade functions: e.g., `execute_swap_base(token_in, token_out, amount)` using Web3.py on Base, and `execute_swap_solana(market, side, amount)` for Solana (possibly integrating Serum DEX via `solana-py` or using Helius for simulation). Test these with small dummy trades on testnet or using a local fork.
- Integrate Helius: call `get_balances` for the Solana wallet, ensure it returns expected (GitHub - vmpyre/helius_sdk: Python SDK for Helius Solana APIs)^{+L253-L260}. Set up a webhook on Helius for a certain program (like a stablecoin mint address) to test receiving webhooks.
- Implement a simple strategy loop that fetches prices (maybe via a public API or by reading on-chain orderbooks) and executes a dummy trade periodically. This is just to confirm the plumbing (that we can detect condition → call execution → see it happen on chain).
- Build the dynamic strategy scaffolding: create a configuration for risk level, trade frequency, etc., that can be tweaked. At this stage, maybe hard-code a simple condition (like if price goes down 5% then buy, up 5% then sell) to test end-to-end trading logic in a deterministic way. Logging for each trade action.

Phase 3: AI Prediction Model Integration

- Develop or integrate a price prediction model. To start, use a simple approach like a moving average crossover or a logistic regression on recent returns – to have something that outputs a buy/sell signal. Ensure the pipeline (data collection → feature → model predict) runs quickly. You can later plug in a more complex model (train an LSTM on historical data offline and load it for inference).
- Integrate the model into the trading loop: instead of dummy condition, use model's prediction to decide trades. For instance, every hour, run prediction for next hour's price move on key assets. If probability of rise > 0.7, set a flag to go long, etc.
- Test this model integration thoroughly in simulation (perhaps record data and simulate actions to see if logically it works). This step ensures the AI aspect is functioning and can be replaced/upgraded independently.

Phase 4: Twitter Bot Core

- Set up Twitter API client with the bot's account. Write a simple script to post a test tweet (like "Hello, I am Based Cheshh"). Verify it appears on Twitter. Delete it after test if needed.
- Implement basic reply logic: use the filtered stream endpoint to listen for tweets from the target influencers (use their user IDs in the filter). On receiving an event, log it for now. Also implement a mention listener (if the bot is mentioned). This might involve using Twitter's Account Activity API or polling mentions timeline if real-time streaming is restricted.
- Write a function for sending replies and one for sending regular tweets (with media support for images). Use Tweepy's `create_tweet` for text and `Client.upload_media` if an image is attached. Test posting a tweet with an image (possibly a meme from a local file).
- Integrate a simple AI text generation for tweets: perhaps use OpenAI's API with a basic prompt ("Write a tweet about Bitcoin in a funny tone"). Test that it returns a plausible tweet and post it. Evaluate format (make sure it's not too long for Twitter).
- Implement scheduling: use Python's `schedule` module or an async loop to schedule tweet tasks. For test, schedule tweets every minute (just for dev) and see that it works (and then remove those). Because actual scheduling will be less frequent, ensure the mechanism triggers reliably. Consider timezones and using cron strings or simple intervals.

Phase 5: NVIDIA Orin & Chart Analyzer

- Set up a method to fetch or generate chart images. Perhaps integrate with a chart API (like TradingView screenshot API) or use matplotlib to plot the last N price points for a coin on the fly. The latter can be done by pulling price data from an exchange API or from our stored history, plotting to an image buffer. For dev, do this for a known coin.
- On the Orin Nano, deploy a pre-trained model for pattern recognition. If NVIDIA provides a sample (maybe a Detectron model for shapes, or a custom one, or even use an existing candlestick pattern detection library), load it. Alternatively, use the **NIM microservice** if available: perhaps NVIDIA has an endpoint for image classification we can use out of the box.

- Write a Python function that sends the chart image to the Orin for analysis. If using a local model, this might be an RPC call or simply running a python script on Orin that listens for images (maybe via ZeroMQ or REST) and returns results. If using NIM cloud, it's a `requests.post` to their API.
- Test the pattern detection with a known scenario: e.g., feed a chart that clearly has a "head and shoulders" and see if the system identifies it (if our model is capable of that). Adjust and calibrate as necessary (maybe the model instead outputs some embedding that we need to interpret). This may require training a model if none exists – which is beyond initial scope, so perhaps rely on simpler heuristics first (like reading OHLC data to identify a few known candlestick patterns as a stopgap, then later incorporate full AI pattern recognition).
- Integrate the output of the chart analysis into the strategy. For instance, add an additional check: if model says "pattern is bearish", maybe override the prediction model if it was bullish (to avoid false signals), etc. Essentially fuse the signals.

Phase 6: Web Dashboard Development

- Scaffold a Next.js app with TypeScript (if desired) and set up pages for Dashboard. Design a clean UI layout using a component library (maybe Material-UI or Chakra UI for speed, or custom Tailwind CSS if aiming for custom design).
- Implement Auth0 authentication on the Next.js app. Follow the quickstart to allow login. Lock down the dashboard routes by checking `user` from Auth0. Test login flow locally with Auth0 dev keys.
- Build the main dashboard page. Use Next.js API routes or `getServerSideProps` to fetch initial data (like current balances, last 5 trades, last tweets). Display those in components (table for trades, list for tweets, etc.). Style with attention to readability – maybe color-code profits vs losses, etc.
- Integrate WebSocket for live updates: possibly use a library like `socket.io` on both server and client, or the WebSocket API directly. Set up a small Node server (could be part of Next.js custom server) that subscribes to Redis pubsub. On message, emit to clients. On the client, connect on component

mount and update state accordingly. Test by manually publishing a fake event in Redis and seeing it appear on the UI.

- Add the Chat Terminal UI on the dashboard (or as a separate page). This could be a simple chat bubble interface. Implement a basic text input that sends messages to an API route. The API route will forward to the LangChain agent and stream back the response (you can use Next.js Edge functions or server-sent events for streaming if needed, or just send the full reply). Initially, just echo messages to test the pipeline. Then integrate with the actual agent.
- Ensure the dashboard auto-refreshes certain stats periodically if needed (or make everything event-driven to avoid polling). For metrics that aren't event-driven (like P&L for the day), maybe update every minute via a small interval or have the backend push periodic summary events.

Phase 7: AI Content & Persona Refinement

- By this phase, the skeletal system is running: trades happen, tweets happen, UI shows stuff. Now focus on refining the **quality of AI outputs**. This involves prompt engineering for tweet generation (e.g., create a prompt template that includes a few example memes, to steer the style). Possibly fine-tune a smaller model on a corpus of crypto tweets to have as a fallback local model.
- Train or refine the sentiment analysis. Possibly integrate a specific model (like a HuggingFace sentiment model fine-tuned on financial tweets). Test it on sample tweets to ensure it categorizes sentiment in a way that matches intuition. Use that output to drive content decisions (maybe implement a function `pick_tone()` that uses sentiment data to pick a tone for the next tweet).
- Test the full engagement cycle: simulate an influencer tweet, see that the bot responds within, say, a minute with an apt reply. Evaluate if the content could be improved; iterate on the prompt or logic. Also simulate some negative scenarios (someone insults the bot) to see how it responds – adjust filters or response patterns to stay likable.
- Use the Virtuals GAME SDK to see if it can enhance decision-making. Perhaps let the GAME agent run for a bit in a sandbox to generate some tweets or replies and observe. If it works well, integrate it into the live loop (or use its output as suggestions that the main bot can use). This step might be optional if LangChain+OpenAI suffice, but could be interesting for adding autonomous

planning (like the agent deciding “I should post 3 times today and then engage with X”).

Phase 8: Testing (Simulated and Live)

- Do a comprehensive test of the entire system in a **simulation mode**: perhaps restrict actual trades (or use very small amounts) and have the bot run for a day or two while closely monitoring. Check logs for any errors or odd behavior. Ensure no funds are lost unexpectedly and tweets are appropriate.
- If all looks good, deploy the bot in **live mode** with real parameters (real trading amounts that are acceptable and connect to mainnet, and real tweeting with the official account). Keep an eye on it especially during initial live runs.
- Have a rollback or kill-switch plan: e.g., be ready to revoke API keys or shut down the server if something goes awry (like a bug causing a trade loop). This could even be a command in the chat (“emergency stop”) that you can use.
- Let the bot run and gather data. Use the dashboard to monitor and also manual cross-check (e.g., verify on Etherscan/SolScan that the transactions match what’s shown, verify on Twitter the posts match logs).

Phase 9: Deployment and Scaling

- Deploy the Next.js app (if not already on the DO server) – possibly containerize the whole app (Docker compose with Python app, Node app, Redis) and deploy to DO droplet. Ensure it restarts on reboot, etc. Use a process manager (systemd or PM2) for resilience.
- Consider domain and HTTPS for the dashboard – maybe set up a domain and SSL (DO can provide Let’s Encrypt easily). This is especially important if Auth0 callback requires a certain domain.
- Enable monitoring: use DO’s monitoring or set up simple alerts (like if the bot process stops or if CPU/network usage goes abnormal).
- Document usage: write internal docs on how to update prompts, how to adjust strategies via the chat or config, so that maintaining the bot or onboarding others to help is easier.

Phase 10: Ongoing Improvement

- After deployment, use the analytics to improve the bot. Perhaps do weekly updates to the model or strategy based on what's working. For example, if arbitrage is consistently profitable, allocate more capital there; if a particular type of meme gets huge engagement, focus the AI on that style.
- Keep the software updated: incorporate user feedback (maybe people interact with the bot on Twitter – if they ask for features or if they find something annoying, tweak accordingly).
- Expand functionality: possibly integrate more chains (if desired), or add Discord integration (the GAME SDK suggests Discord/Telegram support could come – the bot could be extended to those platforms as well).

Each phase should be validated before moving to the next, ensuring a solid foundation. By the end of this roadmap, Based Cheshh would be fully operational and (with luck) performing as an autonomous crypto-trading, meme-slinging entity.

Security Considerations

Throughout development and deployment, we maintain strict security practices for Based Cheshh:

- **Secure Storage of Keys:** All API keys (Twitter, OpenAI, etc.) and crypto private keys are stored in encrypted form or environment variables on the server, never exposed in code repos. For instance, the Solana wallet's private key might be stored encrypted with a passphrase; the bot prompts for this on startup (meaning an operator needs to unlock it when rebooting the bot). This prevents a scenario where someone gaining filesystem access can instantly steal keys. Additionally, use least privilege where possible: the Twitter API keys are limited to the specific bot account; the crypto wallets hold only the funds needed for trading (excess profits can be swept periodically to cold storage).
- **Auth0 & Access Control:** Only authenticated users can access the internal dashboard and chat, as enforced by Auth0. We configure Auth0 to allow login for a small set of emails (the team) or protect it with an additional layer (like IP whitelist for the admin interface). This stops random individuals from logging in even if they somehow got an Auth0 invite link. The chat commands that can

affect the bot's behavior are only via this authenticated channel. The production site is served over HTTPS, so credentials and data in transit are encrypted.

- **Permissions & Scopes:** The Twitter API keys are given only the necessary permissions (read and write tweets). For blockchain, the trading wallet keys have full access to funds in that wallet by necessity, but we could implement smart contract-based controls if desired (e.g., have a governance contract that can lock the bot's funds in emergencies – but this is complex and perhaps overkill initially). If using any cloud services (like if Orin was remote, or any external DB), ensure firewall rules or credentials protect them (e.g., Redis is bound to localhost or requires a password, so it can't be sniffed).
- **Regular Auditing:** We will regularly audit the logs for any unusual behavior. For example, monitor if the bot ever trades outside its intended assets or frequencies – could indicate either a bug or a security issue. Because the bot can execute code (the LLM could theoretically be prompted maliciously to run some unauthorized command if not sandboxed properly), we sanitize inputs. LangChain tools executed by the agent are vetted and cannot run arbitrary OS commands unless explicitly allowed. We do not give the AI agent direct OS-level access beyond its defined tools, preventing injection attacks via the chat or Twitter content.
- **Abuse and Ethical Use:** Since the bot can influence social media, we ensure it doesn't engage in harassment, hate speech, or misinformation. We add content filters for the AI outputs (OpenAI's content filter, or our own keyword checks) to avoid problematic tweets. This is both ethical and to avoid violating Twitter rules. Also, we make it clear in the bio that it's a bot (per Twitter's requirements for automated accounts). This transparency protects us from enforcement actions and builds trust with the community (they know it's an AI gimmick, hopefully endearing rather than deceptive).
- **Scalability & Failover:** While not a direct security issue, having a robust infrastructure means fewer emergencies. We may run the bot on a redundant setup (maybe a secondary instance can take over if primary fails, at least to wind down positions safely). Backing up essential data (like logs, learned model states) so that a server crash doesn't lose all memory of performance.

- **Updating and Maintenance:** We will keep the system updated with security patches. The Next.js app and its dependencies are updated when patches are released (especially for any Auth0 issues or Node vulnerabilities). The Python libs likewise. Docker images (if used) are rebuilt regularly. The NVIDIA Orin system is kept updated (as those often have important security fixes in JetPack). Since the bot deals with finance, any open-source library used for trading (like web3.py, etc.) is monitored for any security advisories.
- **Privacy:** If the bot gathers any user data (currently mainly public Twitter data and maybe Auth0 user emails), we respect privacy. We don't misuse or expose any personal info of the operators or others. This is more of a note for compliance; the main data we have is our own trading and tweeting info.

By considering these points, we aim to prevent both external attacks and internal mistakes. The result is a system that is not only effective but also **trustworthy and secure**, safeguarding funds, data, and reputation. Each part of Based Cheshh is built with a mindset that anticipates potential risks and mitigates them, from the use of proven services like Auth0 to careful handling of keys and limiting the AI's freedoms within safe boundaries.

With these security measures and best practices in place, Based Cheshh can operate continuously and safely, focusing on its core missions of trading and tweeting without compromising on integrity or control.